

Lab 9: Nested and High-Order Functions

(adapted from Profs. Jones and Tolmach's earlier version)

Winter 2015

Nested Functions

- Sometimes it is useful to allow function definitions to be nested, one inside another

```
int f(int x, int y) {  
  int square(int z) { return z*z; }  
  return square(x) + square(y);  
}
```

- This might, for example, be used to introduce a local function, square, without making it more widely visible
- Java, (standard) C, C++ do not allow this, but ML, Pascal, Haskell, and many other languages do ...

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {  
  int[] a;  
  ... a ... readArray ... quicksort ... writeArray ...  
}  
  
void readArray(inp, a) { ... inp ... a ... }  
void writeArray(a, out) { ... a ... out ... }  
  
void quicksort(a, lo, hi) {  
  int pivot = ...;  
  ... a ... pivot ... partition ... quicksort ...  
}  
  
void partition(a, pivot, lo, hi) {  
  ... a ... pivot ... swap ...  
}  
  
void swap(a, i, j) { ... a[i] ... a[j] ... }
```

- Now let's move readArray and writeArray in to sort

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {  
  int[] a;  
  void readArray() { ... inp ... a ... }  
  void writeArray() { ... a ... out ... }  
  
  ... a ... readArray ... quicksort ... writeArray ...  
}  
  
void quicksort(a, lo, hi) {  
  int pivot = ...;  
  ... a ... pivot ... partition ... quicksort ...  
}  
  
void partition(a, pivot, lo, hi) {  
  ... a ... pivot ... swap ...  
}  
  
void swap(a, i, j) { ... a[i] ... a[j] ... }
```

- parameters are no longer required!

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {  
  int[] a;  
  void readArray() { ... inp ... a ... }  
  void writeArray() { ... a ... out ... }  
  
  ... a ... readArray ... quicksort ... writeArray ...  
}  
  
void quicksort(a, lo, hi) {  
  int pivot = ...;  
  void partition(a, pivot, lo, hi) {  
    ... a ... pivot ... swap ...  
  }  
  void swap(a, i, j) { ... a[i] ... a[j] ... }  
  
  ... a ... pivot ... partition ... quicksort ...  
}
```

- Move partition and swap in to quicksort

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {  
  int[] a;  
  void readArray() { ... inp ... a ... }  
  void writeArray() { ... a ... out ... }  
  
  ... a ... readArray ... quicksort ... writeArray ...  
}  
  
void quicksort(a, lo, hi) {  
  int pivot = ...;  
  void partition() {  
    ... a ... pivot ... swap ...  
  }  
  void swap(i, j) { ... a[i] ... a[j] ... }  
  
  ... a ... pivot ... partition ... quicksort ...  
}
```

- again, fewer parameters are required!

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {
    int[] a;
    void readArray() { ... inp ... a ... }
    void writeArray() { ... a ... out ... }

    ... a ... readArray ... quicksort ... writeArray ...
}

void quicksort(a, lo, hi) {
    int pivot = ...;
    void partition() {
        void swap(i, j) { ... a[i] ... a[j] ... }
        ... a ... pivot ... swap ...
    }

    ... a ... pivot ... partition ... quicksort ...
}

... a ... readArray ... quicksort ... writeArray ...
}
```

- Move swap in to partition

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {
    int[] a;
    void readArray() { ... inp ... a ... }
    void writeArray() { ... a ... out ... }

    void quicksort(a, lo, hi) {
        int pivot = ...;
        void partition() {
            void swap(i, j) { ... a[i] ... a[j] ... }
            ... a ... pivot ... swap ...
        }

        ... a ... pivot ... partition ... quicksort ...
    }

    ... a ... readArray ... quicksort ... writeArray ...
}
```

- Move quicksort in to sort

A More Complicated Example: Quicksort

```
void sort(File inp, File out) {
    int[] a;
    void readArray() { ... inp ... a ... }
    void writeArray() { ... a ... out ... }

    void quicksort(lo, hi) {
        int pivot = ...;
        void partition() {
            void swap(i, j) { ... a[i] ... a[j] ... }
            ... a ... pivot ... swap ...
        }

        ... a ... pivot ... partition ... quicksort ...
    }

    ... a ... readArray ... quicksort ... writeArray ...
}
```

- yet again, fewer parameters are required!
- how can we compile code like this?

Exercise

Warning: These exercises require gcc extensions to C that are enabled by default on linux lab machines, but may not be available on Macs (where gcc is not actually really gcc ...)

- Compile and run the program example1.c
- Rewrite it to use nested functions as much as possible, but without changing the parameters to any function. Name this program nested1.c
- Test this new program by compiling and running it
- Now attempt to drop as many parameters as possible from each nested function, but without changing the sequence of calls made. Name this program dropped1.c
- Test this new program by compiling and running it

Implementation: Free Variables


- The challenge here is in dealing with nested functions that access variables that are defined in enclosing functions

```
int f(int x, int y) {
    int g(int z) { return x+z; }
    return g(x+y);
}
```

- For example, x is said to be “bound” in the definition of f, but “free” in the definition of g
- The code for g refers to a variable that is not in its stack frame

Static Links

- To support calls to nested functions like this, we can give the callee a pointer to the stack frame of the “lexically enclosing” function.
- This is known as a *static link* (or *access link*)
- For example, if arguments are passed on stack, we might include the static length as a special “zeroeth” argument



Example

Given the earlier definition:

```
int f(int x, int y) {
    int g(int z) { return x+z; }
    return g(x+y);
}
```

and a call `f(4,2)`, the stack might look something like the following during the call to `g`

In particular, the value for `x` can be found by following the static link `g` and taking the usual offset for `x`

```

y=2
x=4
stack frame for f
static link_f
ret addr_f
dyn link_f
z=6
stack frame for g
static link_g
ret addr_g
dyn link_g
```

Static Link \neq Base Pointer

► Suppose we have the definition:

```
int f(int x, int y) {
    int g(int z) { return h(x+h(z)); }
    int h(int u) { return y*u; }
    return g(x+y);
}
```

► The static link that `g` uses in calls to `h` points to the stack frame for `f`, not the stack frame for `g`

```

y
x
stack frame for f
static link_f
ret addr_f
dyn link_f
z
stack frame for g
static link_g
ret addr_g
dyn link_g
u
stack frame for h
static link_h
ret addr_h
dyn link_h
```

Lexical Depth

- A function at the top level (i.e., with no enclosing function) has lexical depth 0 (and does not need a static link)
- A function that appears inside the definition of a function with depth `n` has depth `n+1`
- To access a variable/call a function at depth `n`, from a function at depth `m` (note that $n \leq m$), then we have to follow the static link in the current frame ($m-n$) times

Lexical Depth for Quicksort

```

/*0*/ void sort(File inp, File out) {
    int[] a;
    /*1*/ void readArray() { ... inp ... a ... }
    /*1*/ void writeArray() { ... a ... out ... }

    /*1*/ void quicksort(a, lo, hi) {
        int pivot = ...;
        /*2*/ void partition() {
            /*3*/ void swap(i, j) { ... a[i] ... a[j] ... }
            ... a ... pivot ... swap ...
        }
        ... a ... pivot ... partition ... quicksort ...
    }
    ... a ... readArray ... quicksort ... writeArray ...
}
```

Static Links in gcc for X86-64

- The X86-64 ABI tries to avoid passing arguments on the stack, and the same philosophy applies to the static link
- So the static link is passed in `%r10` (an otherwise unused caller-save register)
- Even when compiling with optimization on, any variable accessed from a nested function *must* be stored in a stack frame!
 - But still typically no need for `%rbp`; static link points to bottom address of this frame

Worked Example

- Compile the program `example2.c` to a `.s` file using


```
gcc -O1 -S -o example2.s example2.c
```
- Walk step-by-step through the behavior of the assembly code, annotating it
- Show the contents of the stack and key registers at each point

Exercise

Repeat the same steps for `example3.c`

- Annotate the assembly code.
- Identify the lexical depth of each function
- Clearly identify the code that implements a variable reference which spans a lexical depth difference of 2

Higher Order Functions

Functions that take other functions as arguments or return other functions as results are sometimes called *higher order functions*:

```
int(int) f(int x, int y) {  
  int g(int z) { return x+z; }  
  return g;  
}
```

`int(int)` represents a type for functions that take an `int` argument and return an `int` result

- In this case, `g` may be called after `f` has returned
- ... so `g` may need to access `f`'s `x` parameter after `f` has returned
- ... so we can't just dispose of the activation record as soon as a function exits.

Moving from the Stack to the Heap ...

- A solution in this case is to allocate the activation record for `f` on the heap, and not on the stack
- Note that the term "stack frame" is no longer appropriate!
- Unused activation records can be recovered by garbage collection instead of by popping them off the stack
- Alternatively, if we don't keep the stack frames in the heap, then we will need to save a copy of `x`'s value in the representation for `g` before `f` returns
- Variants of these schemes are used in many functional language implementations (and, increasingly now, also in other settings such as C++, Python, and Javascript implementations)

Lambda Expressions

Lambda expressions (anonymous functions) provide a way to write functions without giving them a name.

$x \rightarrow \boxed{\lambda x. e} \rightarrow e$

Haskell	<code>\x -> x + 1</code>
LISP	<code>(lambda (x) (+ x 1))</code>
Python	<code>lambda x: x + 1</code>
Javascript	<code>function (x) x + 1</code>
C++11	<code>[] (int x) -> int { return x + 1; }</code>
Java 8	<code>(int x) -> x + 1</code>

Example

- The previous example:

```
int(int) f(int x, int y) {  
  int g(int z) { return x+z; }  
  return g;  
}
```

- Rewritten using a lambda expression:

```
int(int) f(int x, int y) {  
  return (\z -> x+z);  
}
```

Using Function Values

- A general purpose "mapping" primitive:

```
void mapArray(int(int) f, int[] arr) {  
  for (int i=0; i<arr.length; i++) {  
    arr[i] = f(arr[i]);  
  }  
}
```

Note that the `f` in the body of the `for` loop is a parameter of `mapArray`, not a known function

- To increment every element in an array, `arr`:

```
mapArray(\x -> x + 1, arr);
```

- To double every element in an array, `arr`:

```
mapArray(\x -> x * 2, arr);
```

- Etc...

Composing Function Values

- ▶ A general purpose “composition” primitive:

```
(int)int compose(int(int) f, int(int) g) {
    return \x -> f (g x);
}
```

- ▶ Using `compose`, we can combine two separate mapping operations:

```
mapArray(g, arr);
mapArray(f, arr);
```

- ▶ Into a single iteration across the array:

```
mapArray(compose(f, g), arr);
```

- ▶ Etc...

Representing Function Values

How should we represent values of type `int(int)`?

- ▶ There are many different values, including:
 - $(\lambda z \rightarrow x+z)$, $(\lambda x \rightarrow x+1)$, $(\lambda x \rightarrow x*2)$, $(\lambda x \rightarrow f(g(x)))$, ...
- ▶ ... any of which could be passed as arguments to functions like `mapArray` or `compose`, ...
- ▶ ... so we need a uniform, but flexible way to represent them

A common answer is to represent functions like these by a pointer to a “closure”, a heap allocated record that contains:

- ▶ a code pointer (i.e., the code for the function)
- ▶ the values of its free variables

Because we’re making copies of the free variables, we usually require them to be immutable

Closures

- ▶ Every function of type `int(int)` will be represented using the same basic structure:

codeptr	...
---------	-----

- ▶ The code pointer and list of variables vary from one function value to the next:

$(\lambda z \rightarrow x+z)$	codeptr1	x	
$(\lambda x \rightarrow x+1)$	codeptr2		
$(\lambda x \rightarrow x*2)$	codeptr3		
$(\lambda x \rightarrow f(g(x)))$	codeptr4	f	g

Constructing and Calling Closures

We can encode closure construction in standard IR code:

- ▶ make an RTS call to allocate the closure record
- ▶ store the values of the free variables into the record

To call a function via a closure:

- ▶ add a pointer to the closure as an extra initial argument (to provide access to any free variables)
- ▶ make an indirect call to the code pointed to by the first field of the closure

Within the function code, free variables are referenced via the closure argument

- ▶ Known functions without free variables don’t need a closure argument and can be called directly

Implementation of $(\lambda z \rightarrow x+z)$

```
...
t1 = call _malloc(12)
0[t1]:P = _code1
8[t1]:I = x
...

_code1(clo,z)
(x)
{
    x = 8[clo]:I
    t1 = x+z
    return t1
}
```

Implementation of $(\lambda x \rightarrow x+1)$

```
...
t1 = call _malloc(8)
0[t1]:P = _code2
...

_code2(clo,x)
{
    t1 = x+1
    return t1
}
```

Implementation of $(\lambda x \rightarrow f(g(x)))$

```
...
t1 = call _malloc(24)
0[t1]:P = _code4
8[t1]:P = f
16[t1]:P = g
...

_code4(clo,x)
(f,g)
{
  f = 8[clo]:P
  g = 16[clo]:P
  t1 = 0[g]:P
  t2 = call *t1(g,x)
  t3 = 0[f]:P
  t4 = call *t3(f,t2)
  return t4
}
```

Exercises

- Fill in the code for these three functions in `example4.ir`. For simplicity, assume that they are always called directly, rather than via closures, so they have no closure argument

```
(int)int compose(int(int) f, int(int) g) {
  return \x -> f (g x);
}
(int)int h(int x) {
  return compose(\z -> x + z,
                \x -> x + 1);
}
int top() {
  return (h(42))(0);
}
```

- Test your code using the IR interpreter

Closures vs. Objects

Invoking an unknown function through a closure is very similar to invoking a method of an object ...

- Recall that method invocations pass the object itself as an implicit argument
- Closures are like objects with a single method
- Free variables correspond to object fields

In Java 8, lambda expressions are “just” a convenient way to write (local, anonymous) definitions of single-method classes

Very useful for GUI call-backs, aggregate operations, etc.

Summary

- More sophisticated forms of function can be supported by modifying or generalizing how activation records are used
- For a language with higher-order functions, we need to allocate activation records on the heap, or copy data to other heapallocated objects, because a function value may have a longer lifetime than the function that created it
- Function values can be represented by closure records that pair a code pointer with a list of variable values
- Invoking an unknown function through a closure is very similar to invoking a method of an object ...
- The techniques shown here are key tools in the implementation of functional programming languages and are gradually becoming common in OO languages too